

Practical 7

21-2-25

Sharding using Mongoddb

Write-up: -

- Sharding
- Sharding Architecture
- Sharding Types, benefits and limitations

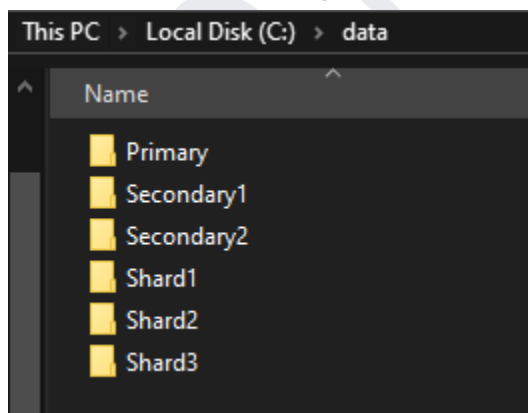
You are a database administrator tasked with setting up MongoDB sharding to handle a growing dataset. You need to configure a MongoDB sharded cluster with the following requirements:

- One Config Server (Replica set)
- Two Shard Servers
- One Mongos Router
- Sharding enabled for a database and collection

<https://gist.github.com/DGamer007/0864c6aeebf27e3821602d9dd5ca737>

5

Create the following folders



Creating 3 Replica Servers

In 3 powershells:

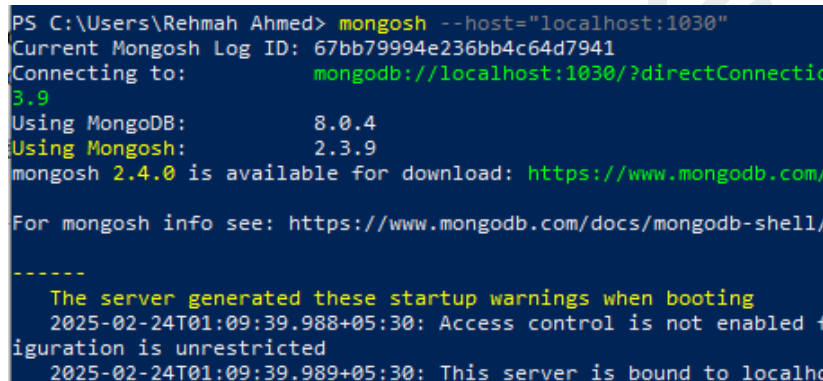
```
mongod --configsvr --port=1030 --replSet="Prac7Repl"  
--dbpath="E:\data\Primary"
```

```
mongod --configsvr --port=1040 --replSet="Prac7Repl"  
--dbpath="E:\data\Secondary1"
```

```
mongod --configsvr --port=1050 --replSet="Prac7Repl"  
--dbpath="E:\data\Secondary2"
```

Connect to anyone of them using mongosh & Initiate Replica Set In a new powershell

```
mongosh --host="localhost:1030"
```



```
PS C:\Users\Rehmah Ahmed> mongosh --host="localhost:1030"  
Current Mongosh Log ID: 67bb79994e236bb4c64d7941  
Connecting to:      mongodb://localhost:1030/?directConnectio  
3.9  
Using MongoDB:      8.0.4  
Using Mongosh:      2.3.9  
mongosh 2.4.0 is available for download: https://www.mongodb.com/  
For mongosh info see: https://www.mongodb.com/docs/mongodb-shell/  
  
-----  
The server generated these startup warnings when booting  
2025-02-24T01:09:39.988+05:30: Access control is not enabled +  
figuration is unrestricted  
2025-02-24T01:09:39.989+05:30: This server is bound to localh
```

```
rs.initiate({  
  _id:"Prac7Repl",  
  configsvr:true,  
  members:[  
    {_id:0, host:"localhost:1030"},  
    {_id:1, host:"localhost:1040"},  
    {_id:2, host:"localhost:1050"}  
  ]  
})
```

```
test> rs.initiate({
...   _id:"Prac7Repl",
...   configsvr:true,
...   members:[
...     {_id:0, host:"localhost:1030"},
...     {_id:1, host:"localhost:1040"},
...     {_id:2, host:"localhost:1050"}
...   ]
... })
{
  ok: 1,
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1740339617, i: 1 }),
    signature: {
      hash: Binary.createFromBase64('AAAAAAAAAAAAAAAAAAAAAAAAAAAA=', 0),
      keyId: Long('0')
    }
  },
  operationTime: Timestamp({ t: 1740339617, i: 1 })
}
```

Creating 3 Shard Instances

In 3 powershells:

```
mongod --shardsvr --port=1130 --dbpath="E:\data\Shard1"
--replSet="Prac7Shard"
```

```
mongod --shardsvr --port=1140 --dbpath="E:\data\Shard2"
--replSet="Prac7Shard"
```

```
mongod --shardsvr --port=1150 --dbpath="E:\data\Shard3"
--replSet="Prac7Shard"
```

Connect to anyone of them using mongosh and Initiate Shard

In a new powershell

```
mongosh --host="localhost:1130"
```

```
PS C:\Users\Rehmah Ahmed> mongosh --host="localhost:1130"
Current Mongosh Log ID: 67bb79e372472e08be4d7941
Connecting to:      mongodb://localhost:1130/?directConnection=true&serverSelectionTime
3.9
Using MongoDB:      8.0.4
Using Mongosh:      2.3.9
mongosh 2.4.0 is available for download: https://www.mongodb.com/try/download/shell

For mongosh info see: https://www.mongodb.com/docs/mongodb-shell/

-----
The server generated these startup warnings when booting
2025-02-24T01:10:41.654+05:30: Access control is not enabled for the database. Read and
figuration is unrestricted
2025-02-24T01:10:41.657+05:30: This server is bound to localhost. Remote systems will be
server. Start the server with --bind_ip <address> to specify which IP addresses it should se
bind_ip_all to bind to all interfaces. If this behavior is desired, start the server with
this warning
-----
```

```
rs.initiate({
  _id:"Prac7Shard",
  members: [
    { _id: 0, host: "localhost:1130"},
    { _id: 1, host: "localhost:1140"},
    { _id: 2, host: "localhost:1150"}
  ]
})
```

```
test> rs.initiate({
...   _id:"Prac7Shard",
...   members: [
...     { _id: 0, host: "localhost:1130"},
...     { _id: 1, host: "localhost:1140"},
...     { _id: 2, host: "localhost:1150"}
...   ]
... })

{
  ok: 1,
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1740339687, i: 1 }),
    signature: {
      hash: Binary.createFromBase64('AAAAAAAAAAAAAAAAAAAAAAAAAAAA=', 0),
      keyId: Long('0')
    }
  },
  operationTime: Timestamp({ t: 1740339687, i: 1 })
}
```

Initialize a Query Router which is a mongos process.

mongos --port=1210

--configdb="Prac7Repl/localhost:1030,localhost:1040,localhost:1050"

```
{
  "t": {"$date": "2025-02-24T01:28:26.119+05:30"},
  "s": "I",
  "c": "COMMAND",
  "id": 51803,
  "svc": "R",
  "ctx": "ClusterS
parameterRefresher",
  "msg": "Slow query",
  "attr": {
    "type": "command",
    "isFromUserConnection": false,
    "ns": "admin.$cmd",
    "col
Type": "admin",
    "command": {
      "commitTransaction": 1,
      "writeConcern": {
        "w": "majority",
        "wtimeout": 0,
        "lsid": {
          "id": {
            "$uuid"
7db-23c0-4ffd-9b22-f7c6dbcb209e"
        },
        "uid": {
          "$binary": {
            "base64": "47DEQpj8HBSa+/TImW+5JCeuQeRkm5NMpJWZG3hSuFU="
          },
          "subT
"
        },
        "txnNumber": 33,
        "autocommit": false,
        "$db": "admin",
        "numYields": 0,
        "reslen": 183,
        "locks": {},
        "protocol": "op_msg",
        "opWaitMillis": 153,
        "queues": {
          "execution": {},
          "ingress": {},
          "workingMillis": 154,
          "durationMillis": 154
        }
      }
    },
    "t": {
      "$date": "2025-02-24T01:28:30.998+05:30"},
    "s": "I",
    "c": "CONNPOOL",
    "id": 22572,
    "svc": "-",
    "ctx": "ShardReg
"msg": "Dropping all pooled connections",
    "attr": {
      "hostAndPort": "localhost:1130",
      "error": "ConnectionPoolExpired: Po
localhost:1130 has expired."
    }
  },
  "t": {
    "$date": "2025-02-24T01:28:56.061+05:30"},
    "s": "I",
    "c": "TXN",
    "id": 51805,
    "svc": "R",
    "ctx": "ClusterS
parameterRefresher",
    "msg": "transaction",
    "attr": {
      "parameters": {
        "lsid": {
          "id": {
            "$uuid": "16d0b7db-23c0-4ffd-9b22-f7c6db
",
            "uid": {
              "$binary": {
                "base64": "47DEQpj8HBSa+/TImW+5JCeuQeRkm5NMpJWZG3hSuFU="
              },
              "subType": "0"
            }
          },
          "txnNumber": 34,
          "txnR
nter": 0,
          "autocommit": false,
          "readConcern": {
            "level": "snapshot"
          },
          "globalReadTimestamp": {
            "ts": Timestamp(1740340734,
            "numParticipants": 1,
            "terminationCause": "committed",
            "commitType": "singleShard",
            "commitDurationMicros": 95297,
            "timeA
icros": 99486,
            "timeInactiveMicros": 745,
            "durationMillis": 100
          }
        },
        "t": {
          "$date": "2025-02-24T01:30:37.789+05:30"},
          "s": "I",
          "c": "NETWORK",
          "id": 6496702,
          "svc": "-",
          "ctx": "ReplicaS
or-TaskExecutor",
          "msg": "Acquired connection for remote operation and completed writing to wire",
          "attr": {
            "duration
```

Now, Connect Shards and Query Router (mongos)

mongosh --host="localhost:1210"

```
PS C:\Users\Rehmah Ahmed> mongosh --host="localhost:1210"
Current Mongosh Log ID: 67bb7c9ebc4c0639264d7941
Connecting to: mongod://localhost:1210/?directConnection=true&ser
Using MongoDB: 8.0.4
Using Mongosh: 2.3.9
mongosh 2.4.0 is available for download: https://www.mongodb.com/try/downlo
For mongosh info see: https://www.mongodb.com/docs/mongod-shell/

-----
The server generated these startup warnings when booting
2025-02-24T01:11:54.701+05:30: Access control is not enabled for the data
2025-02-24T01:11:54.702+05:30: This server is bound to localhost. Remote
esses it should serve responses from, or with --bind_ip_all to bind to all
-----
```

sh.addShard("Prac7Shard/localhost:1130,localhost:1140,localhost:1150")

```
[direct: mongos] test> sh.addShard("Prac7Shard/localhost:1130,localhost:
{
  shardAdded: 'Prac7Shard',
  ok: 1,
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1740340385, i: 1 }),
    signature: {
      hash: Binary.createFromBase64('AAAAAAAAAAAAAAAAAAAAAAAAAAAA= ', 0),
      keyId: Long('0')
    }
  },
  operationTime: Timestamp({ t: 1740340385, i: 1 })
}
```

sh.enableSharding("practice")

```
[direct: mongos] test> sh.enableSharding("practice")
{
  ok: 1,
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1740340390, i: 1 }),
    signature: {
      hash: Binary.createFromBase64('AAAAAAAAAAAAAAAAAAAAAAAAAAAA=', 0),
      keyId: Long('0')
    }
  },
  operationTime: Timestamp({ t: 1740340390, i: 1 })
}
```

Shard a Collection on the Sharding Enabled Database

sh.shardCollection("practice.students", { "enroll": "hashed" })

```
[direct: mongos] test> sh.shardCollection("practice.students", { "enroll": "hashed" })
{
  collectionssharded: 'practice.students',
  ok: 1,
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1740340403, i: 3 }),
    signature: {
      hash: Binary.createFromBase64('AAAAAAAAAAAAAAAAAAAAAAAAAAAA=', 0),
      keyId: Long('0')
    }
  },
  operationTime: Timestamp({ t: 1740340403, i: 3 })
}
```

Inserting 6 values

use practice

```
db.students.insertMany([
  { _id: 1, name: "Alice", enroll: 1001 },
  { _id: 2, name: "Bob", enroll: 2005 },
  { _id: 3, name: "Charlie", enroll: 1503 },
  { _id: 4, name: "David", enroll: 2507 },
  { _id: 5, name: "Eve", enroll: 3002 },
  { _id: 6, name: "Uno", enroll: 3501 }
])
```

```
...    { _id: 1, name: "Alice", enroll: 1001 },
...    { _id: 2, name: "Bob", enroll: 2005 },
...    { _id: 3, name: "Charlie", enroll: 1503 },
...    { _id: 4, name: "David", enroll: 2507 },
...    { _id: 5, name: "Eve", enroll: 3002 },
...    { _id: 6, name: "Uno", enroll: 3501 }
...  ])
{
  acknowledged: true,
  insertedIds: { '0': 1, '1': 2, '2': 3, '3': 4, '4': 5, '5': 6 }
}
```

Checking how they are stored

sh.status()

```
[direct: mongos] practice> sh.status()
shardingVersion
{ _id: 1, clusterId: ObjectId('67bb79af06c9b7dedf09a0e1') }
---
shards
[
  {
    _id: 'Prac7Shard',
    host: 'Prac7Shard/localhost:1130,localhost:1140,localhost:1150',
    state: 1,
    topologyTime: Timestamp({ t: 1740339736, i: 2 }),
    replSetConfigVersion: Long('-1')
  }
]
---
active mongoses
[ { '8.0.4' : 1 } ]
---
autosplit
{ 'Currently enabled': 'yes' }
---
balancer
{
  'Currently running': 'no',
  'Currently enabled': 'yes',
  'Failed balancer rounds in last 5 attempts': 0,
  'Migration Results for the last 24 hours': 'No recent migrations'
}
---
shardedDataDistribution
[
  {
    ns: 'practice.students',
    shards: [
      {
        shardName: 'Prac7Shard',
        numOrphanedDocs: 0,
        numOwnedDocuments: 6,
        ownedSizeBytes: 246,
        orphanedSizeBytes: 0
      }
    ]
  }
],
```

Find By Enroll (Efficient)

db.students.find({ enroll: 2005 })

```
[direct: mongos] practice> db.students.find({ enroll: 2005 })  
[ { _id: 2, name: 'Bob', enroll: 2005 } ]
```

Range queries (not efficient)

db.students.find({ enroll: { \$gt: 1000, \$lt: 3000 } })

MongoDB has to check all shards since hashed values are scattered.